

A Review of: Generalized Cache Tiling for Dataflow Programs

Zachary Sisco

April 12, 2017

Overview

“Generalized Cache Tiling for Dataflow Programs” by Łukasz Domagała, Duco van Amstel, and Fabrice Rastello. LCTES 2016.

- ▶ What is dataflow programming? Why?
- ▶ Extending tiling to dataflow programs
- ▶ Cache-tiling optimization solutions
- ▶ Experimental Results & Analysis
- ▶ Conclusions & Future Work

Dataflow Programming

Controlflow Program

- ▶ Stream of instructions operate on external data.
- ▶ Conditional execution change the instruction-execution path.
- ▶ Data is “static” unless instruction stream moves it.

Dataflow Program

- ▶ Stream of data passing from instruction to instruction.
- ▶ Conditional execution routes data to different instructions.
- ▶ Data flowing through “static” instructions.

Dataflow Programming

Applications

- ▶ Data streaming and processing (e.g., video processing)
- ▶ Cryptographic encoding and decoding
- ▶ Embedded systems

Loop Tiling

Group together operations within an iteration space so they execute atomically.

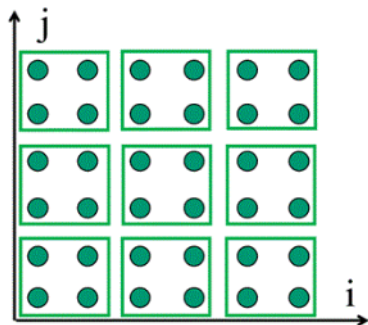


Figure 1: Reduce cache misses with loop-nest tiling. Figure from (Liu, 2017)

Dataflow Iteration Space

- ▶ Actors — basic programmable units
- ▶ Channels — connect actors, pass data elements

Dataflow Iteration Space

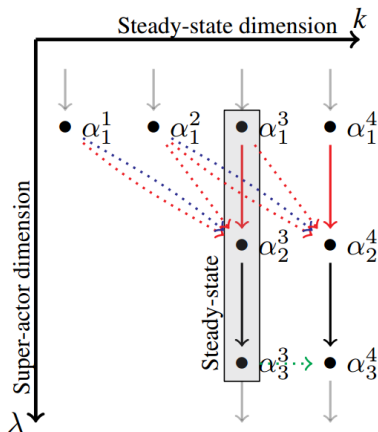


Figure 2: Iteration space with dependencies. From (Domagała et al., 2016)

Dataflow Iteration Space

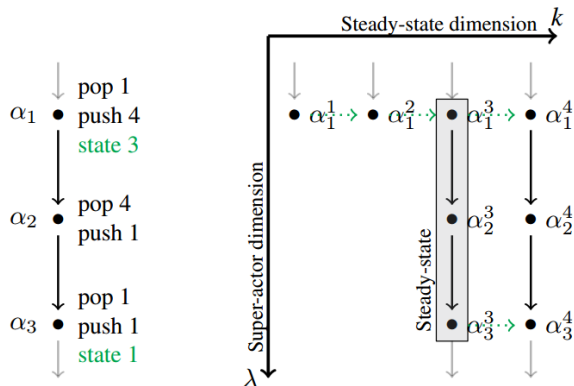


Figure 3: Iteration space after transformation. From (Domagała et al., 2016)

Optimization Problem

Minimize cache misses

- ▶ Find the largest possible width for each tile that satisfies its cache-size requirement and minimizes a cost function.
- ▶ Cost function is defined with respect to minimizing the variable *spill*.

Constraint Programming

Constraints

- ▶ Ensure the ordering of nodes respects the original dataflow dependencies
- ▶ Ensuring the unique indexing and ordering of nodes
- ▶ Ensuring the cache size is not exceeded
- ▶ And more...

Heuristics-based Solutions

Near-optimal solution in linear-time complexity

- ▶ Greedy tiling
- ▶ Schedule actors based on descending order of edge weights (size of data transfer)

Experimental Evaluation

- ▶ StreamIt dataflow language
- ▶ Benchmark of 12 dataflow programs
- ▶ Simulated ARM architecture with 16kB L1 instruction cache and 8kB L1 data cache
- ▶ Least-Recently Used cache replacement policy

Experimental Evaluation

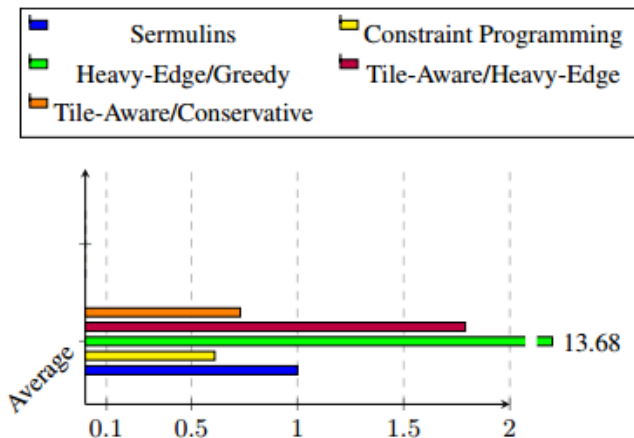


Figure 4: Average cache-miss amounts. Adapted from (Domagała et al., 2016)

Conclusions & Future Work

- ▶ Reduce cache misses in dataflow programs through cache-tiling optimization
- ▶ Improves over existing optimizations
- ▶ Extension to parallel or distributed architectures
- ▶ Improve model with information about cache layout/associativity

Thank You

Questions?